# UNDERSTANDING WORKING ON SNOWFLAKE UNISTORE ZERO TO HERO

#### **BASIC DISCUSSION**

#### 1. What is a Transactional (OLTP) workload?

Explanation: Transactional systems are optimized for frequent, short, real-time operations like inserts, updates, and deletes.

**finTech Use Case:** In **Revolut**, each user tap (adding money, transferring funds) is a transaction that must be processed in **milliseconds** — no delay allowed.

#### **X** Example (OLTP-style query):

-- Inserting a payment transaction INSERT INTO transactions (transaction\_id, user\_id, amount, type, timestamp) VALUES ('txn\_1001', 'usr\_001', 250.00, 'DEBIT', CURRENT\_TIMESTAMP);

#### 2. What is an Analytical (OLAP) workload?

Explanation:OLAP systems are built for reading large datasets, doing aggregations, and uncovering trends — typically not for real-time updates.

**Use Case:Netflix** analyzes watch history across millions of users to personalize content. This doesn't need millisecond inserts, but rather powerful scanning and aggregations.

#### **X** Example (OLAP-style query):

-- Analyze monthly total spending per user SELECT user\_id, MONTH(timestamp) AS txn\_month, SUM(amount) AS total\_spend FROM transactions GROUP BY user\_id, MONTH(timestamp);

#### 3. What is a Primary Key, and why is it a real deal in Hybrid tables?

- Explanation: A Primary Key is a column (or set of columns) that uniquely identifies every row in a table no duplicates, no NULLs allowed.
- FinTech Example: Stripe must ensure that every transaction has a unique transaction\_id it cannot afford to lose or overwrite records.

#### **X** SQL Example:

```
CREATE OR REPLACE HYBRID TABLE transactions (
transaction_id STRING PRIMARY KEY,
user_id STRING,
amount NUMBER,
timestamp TIMESTAMP
);
```

### 4. What is a Composite Primary Key and when is it useful?

- Explanation: A Composite Key combines two or more columns to uniquely identify a row. Useful when no single field is unique on its own.
- **1** Use Case:In Plaid, a user\_id and bank\_id might form a composite key to uniquely identify a bank account across users and institutions.

#### **SQL** Example:

```
CREATE OR REPLACE HYBRID TABLE user_bank_accounts (
    user_id STRING,
    bank_id STRING,
    account_type STRING,
    PRIMARY KEY (user_id, bank_id)
);
```

Composite keys prevent duplicate account entries for the same user and bank.

#### 5. What is a Foreign Key and how does it enforce relationships?

Explanation: A Foreign Key links one table to another via a primary key, ensuring data integrity. It enforces that referenced data must exist.

FinTech Example: In Robinhood, each trade must refer to an existing user — a foreign key enforces this.

#### **X** SQL Example:

```
CREATE OR REPLACE HYBRID TABLE users (
user_id STRING PRIMARY KEY,
name STRING
);

CREATE OR REPLACE HYBRID TABLE trades (
trade_id STRING PRIMARY KEY,
user_id STRING,
stock_symbol STRING,
trade_amount NUMBER,
FOREIGN KEY (user_id) REFERENCES users(user_id)
);
```

### 6. What is Referential Integrity and why is it critical?

**Explanation:Referential Integrity** ensures that the connection between tables is valid — no orphan records. You can't add a trade if the user doesn't exist.

#### Nalid Insert:

-- This will FAIL if 'usr\_999' does not exist in users table INSERT INTO trades VALUES ('trade\_001', 'usr\_999', 'AAPL', 1000);

This protects your data quality — crucial for auditability in finance.

#### 7. What is an Index and how does it speed up queries?

**Explanation:**An **index** acts like a search engine. Without one, the system scans the entire table. With one, it goes straight to the data.

**TECH-USE CASE:Amazon** needs to retrieve orders by user\_id instantly. Indexes make this fast.

#### **X SQL Example:**

-- Create index on user\_id for quick filtering CREATE INDEX idx user id ON transactions(user id);

# 8. What's the problem with maintaining separate OLTP and OLAP systems?

- Explanation: You need to build pipelines to move data between them it causes:
  - Latency in insights
  - Duplicated infrastructure
  - Security headaches

Real-world Example: Goldman Sachs might move data from MongoDB (for transactions) to Snowflake (for analytics), delaying fraud detection. Unistore removes this barrier.

#### 9. Why does real-time access matter in FinTech?

**Explanation:** Decisions like fraud detection or credit scoring must happen **now**, not later.

**Example:**If **Zelle** delays detection, fraudulent transactions may be processed before detection kicks in. Unistore enables **immediate analytical insights** on real-time transactional data.

## 10. How does Snowflake unify OLTP and OLAP with Unistore?

■ Explanation:Unistore uses Hybrid Tables backed by a row-based engine for fast OLTP-like operations, but still seamlessly integrates with columnar storage for OLAP analytics.

#### TARCHITECTURE Insight:

- Hybrid Tables: Row-based, fast inserts/selects
- Analytical Store: Columnar, scalable queries
- Shared Engine: Same SQL, auth, metadata, compute
- Invisible Replication: No manual sync needed

# Part 2: Deeper DISCUSSIONS About Unistore & Hybrid Tables (ETL, Architecture, Use)

#### 1. How do Hybrid Tables work internally in Snowflake?

- Explanation: Hybrid Tables:
  - Use **row-based storage** → Fast for single inserts/updates
  - Are replicated automatically to the analytic store
  - Leverage same compute/query engine as Snowflake
- Difference: Unlike traditional Snowflake tables, Hybrid Tables don't use micro-partitions (columnar), but row storage for OLTP efficiency.
- @ Benefit: You don't need to choose "which engine" it's automatically managed.

#### 2. How can I build an end-to-end ETL/ELT pipeline with Unistore?

#### Steps:

- Load Raw Data into Hybrid Tables (e.g., payments, accounts)
- Transform with SQL (joins, filters, aggregations)
- Store or serve insights in regular Snowflake tables or dashboards
- Feed back transformed data to front-end apps in milliseconds

#### X SQL Example:

- -- Load transactions INSERT INTO hybrid transactions VALUES (...);
- -- Transform: Aggregate user daily spend CREATE OR REPLACE TABLE daily\_summary AS SELECT user\_id, DATE(timestamp) AS txn\_date, SUM(amount) AS daily\_spend FROM hybrid\_transactions GROUP BY user\_id, DATE(timestamp);

# 3. What use cases does Unistore unlock that weren't possible before?

- Real-time fraud detection using current + historical data
- In-app analytics dashboards powered directly from transactional tables
- Alerting systems reacting instantly to operational data

**Example:SoFi** could detect a credit threshold breach and instantly trigger UI alerts + credit score checks — all using Unistore without syncing across DBs.

# 4. Can Hybrid Tables be used in joins with regular Snowflake tables?

Yes!Snowflake allows joins between Hybrid and regular tables, treating them as a single logical dataset.

#### **X** SQL Example:

SELECT t.user\_id, u.name, t.amount FROM hybrid\_transactions t JOIN users u ON t.user id = u.user id;

# 5. How does Unistore simplify app development in FinTech or TECH environments?

**Explanation:** Developers can now:

- Store real-time app data in Hybrid Tables
- Analyze instantly in dashboards
- Trigger alerts or update workflows
- ullet Eliminate syncing layers like Kafka o ETL o Snowflake
- ${\color{red} \overline{\triangleright}}$  Before Unistore:Frontend App  $\rightarrow$  OLTP DB  $\rightarrow$  ETL  $\rightarrow$  Snowflake  $\rightarrow$  BI Tools
- With Unistore: Frontend App ↔ Snowflake only (via Hybrid Tables)
- **Example:Uber's** driver payment app could run entirely on Snowflake, storing trips as Hybrid Table records and instantly calculating payouts.

## Use Case: Real-time Credit Breach Detection & Alert System in PayPal using Snowflake Unistore

- Business Scenario PayPal wants to:
  - Track customer credit usage in real time
  - Instantly alert users when they exceed their allowed credit
  - Provide dashboards to analysts showing breach trends and usage patterns

#### This must be built on **one system** without using:

- An external OLTP DB (like PostgreSQL)
- Middleware queues (Kafka)
- Separate ETL tools

#### **Unistore + Hybrid Tables** in Snowflake makes this possible.

# T System Architecture Overview

#### Components:

- 1. customers → Static metadata table (regular Snowflake table)
- 2. hybrid\_transactions → Live credit transactions (Hybrid Table)
- 3. credit\_breach\_alerts → Alert log (Hybrid Table)
- 4. daily\_summary → Aggregated analytics (Regular Table or Materialized View)
- 5. **SQL logic / tasks / streams** → For alerting and transformation

# Step 1: Set up base tables

```
A. Customer Master Table (Regular Table)
CREATE OR REPLACE TABLE customers (
 customer id STRING PRIMARY KEY,
 name STRING.
credit limit NUMBER
);
-- Sample data
INSERT INTO customers VALUES
 ('CUST001', 'Alice', 1000),
 ('CUST002', 'Bob', 500),
 ('CUST003', 'Charlie', 750);
B. Transactions Table (HYBRID TABLE)
CREATE OR REPLACE HYBRID TABLE hybrid transactions (
 txn id STRING PRIMARY KEY,
 customer id STRING.
 amount NUMBER,
 txn time TIMESTAMP.
 FOREIGN KEY (customer_id) REFERENCES customers(customer_id)
);
-- Simulate incoming transactions
INSERT INTO hybrid transactions VALUES
 ('TXN1001', 'CUST001', 300, CURRENT TIMESTAMP),
 ('TXN1002', 'CUST001', 400, CURRENT TIMESTAMP),
 ('TXN1003', 'CUST002', 600, CURRENT TIMESTAMP); -- breach
```

Referential Integrity is enforced: no transaction for non-existing customers

## Step 2: Alert Table for Credit Breach (HYBRID TABLE)

```
CREATE OR REPLACE HYBRID TABLE credit_breach_alerts (
    alert_id STRING PRIMARY KEY,
    customer_id STRING,
    total_spent NUMBER,
    credit_limit NUMBER,
    breach_time TIMESTAMP
);
```

#### Step 3: Create logic to detect breach (ELT with INSERT-SELECT)

This logic compares a customer's total spend in the hybrid\_transactions table against their credit limit in customers.

```
-- Insert into alerts table if spending > credit limit
INSERT INTO credit_breach_alerts
SELECT
UUID_STRING() AS alert_id,
t.customer_id,
SUM(t.amount) AS total_spent,
c.credit_limit,
CURRENT_TIMESTAMP
FROM hybrid_transactions t
JOIN customers c ON t.customer_id = c.customer_id
GROUP BY t.customer_id, c.credit_limit
HAVING SUM(t.amount) > c.credit_limit;
```

Step 4: Analytical Summary for Dashboard: This could go to a dashboard like Looker/Tableau/Streamlit.

```
-- Daily spend summary for internal insights
CREATE OR REPLACE TABLE daily_summary AS
SELECT
customer_id,
DATE(txn_time) AS txn_date,
SUM(amount) AS daily_spend,
COUNT(*) AS txn_count
FROM hybrid_transactions
GROUP BY customer_id, DATE(txn_time);
```

# Optional: Automate using TASKS

To auto-run the alert logic every minute (real-time reaction):

CREATE OR REPLACE TASK detect\_credit\_breach
WAREHOUSE = compute\_wh
SCHEDULE = '1 MINUTE'
AS
INSERT INTO credit\_breach\_alerts
SELECT
 UUID\_STRING(), t.customer\_id, SUM(t.amount), c.credit\_limit,
CURRENT\_TIMESTAMP
FROM hybrid\_transactions t
JOIN customers c ON t.customer\_id = c.customer\_id
GROUP BY t.customer\_id, c.credit\_limit
HAVING SUM(t.amount) > c.credit\_limit;

# Outcome of the Pipeline

| Component | гуре          | Function                   |
|-----------|---------------|----------------------------|
| m o u o   | Degular table | Holds motodata like credit |

| customers            | Regular table    | Holds metadata, like credit limit   |
|----------------------|------------------|-------------------------------------|
| hybrid_transactions  | Hybrid Table     | Stores real-time transactional data |
| credit_breach_alerts | Hybrid Table     | Triggers & logs alerts immediately  |
| daily_summary        | Analytical table | Power dashboard metrics             |
| detect_credit_breach | Task             | Automates detection every minute    |

# Real-World Benefits for PayPal

| Feature               | Benefit                                       |
|-----------------------|---|
| Phybrid Tables        | Fast inserts for real-time processing         |
| Referential Integrity | Guarantees data quality (no orphan txns)      |
| In Unified Store      | No separate ETL or DB needed                  |
| Automation            | Auto-monitor credit and send alerts instantly |
| ✓ Instant Analytics   | Finance & Risk teams see live trends          |

# What you learnt in this Lab: If you're onboarding a FinTech Data Engineer, this lab:

- Teaches OLTP + OLAP unification
- Introduces real-time alerting logic
- Uses just Snowflake SQL (no third-party tools)
- Is extendable: Add email alerts via Snowflake External Functions or Streamlit UIs for visual alerts

Happy learning Best regards Saransh Jain